Communication in Electronic Systems

Lecture 8: Networking and Transport Layers

Lecturer: Petar Popovski

TA: Junya Shiraishi, João H. Inacio de Souza

email: petarp@.es.aau.dk





Course Overview: Part 2. Communication and Networking

- MM5: Introduction to Communication Systems
- MM6: Simple Multiuser Systems and Layered System Design
- MM7: Network Topology and Architecture
- MM8: Networking and Transport Layers
- MM9: Introduction to Security
- Guest lecture
- MM10: Packets and Digital Modulation
- MM11: Communication Waveforms
- MM12: Workshop on Modulation and Link Operation

outline

- network layer
- routing
- transport layer
- TCP connection
- congestion control
- fairness



recap: lecture 3

lecture 3:
protocols and network infrastructure

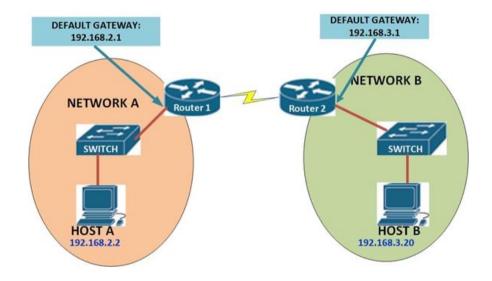


- why is layering a good architectural solution for protocols?
- how to ease the implementation of the network?
 - o idea of communication modules and layering
- why do we need networking?
- network systems and topologies have been discussed



network layer

- transport segment from sending to receiving host
 - on sending side:encapsulates segments into datagrams
 - on receiving side:
 delivers segments to transport layer
 - o principle of store-and-forward
- network layer protocols in every host, router
- the lowest layer that deals with end-to-end transmission



source: https://www.ccnablog.com/network-layer/



functionality of the network layer (1)

- interfacing to the transport layer and the data link layer (OSI) / the host-to-network layer (TCP/IP).
- goals:
 - services independent of router technology
 - transport layer shielded from the number, type, and topology of routers
 - network addresses with uniform numbering/identification plan
- the most important functionality is routing:
 - make sure that the packets can find their way through the networks.
 - QoS (Quality of Service) can also be handled here, and so can some elements of congestion control.
- analogy: the (snail) mail system.
 - a mail is sent by putting it in the mail box, in the required format
 - the mail system takes care that it will arrive at the receiver



Transport

Network

DLC layer

link 2

DLC layer

link 3

laver

DLC layer

link 1

layer





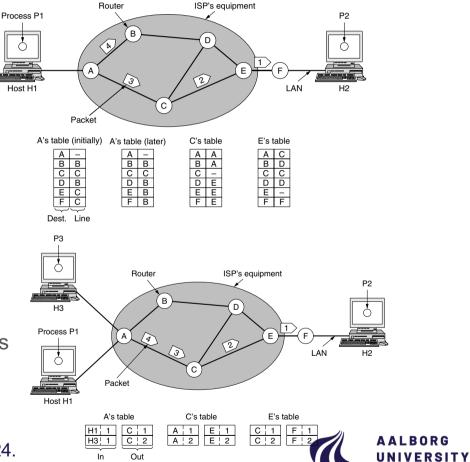
connectionless vs. connection-oriented services

connectionless

- modelled by the postal service
- no prior setup
- no readiness to receive,e.g., random access
- suitable for multicast

connection-oriented

- modelled by the phone service
- initial handshake to establish
- low control information/metadata afterwards
- suitable for longer, frequent exchanges



Petar Popovski, Communication in Electronic Systems, Fall 2024.

functionality of the network layer (2)

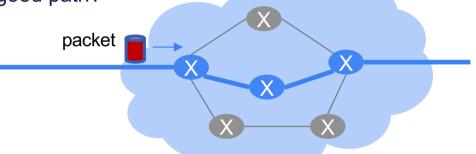
- the network layer can offer either connectionless or connection-oriented services to the transport layer:
 - o ATM (Asynchronous Transfer Mode): Connection-oriented with virtual circuits
 - IP: Connectionless, with datagrams (store and forward)
- even if the network layer is connectionless,
 the transport layer can still be connection-oriented (e.g. TCP).
- the internet camp thinks network layer should do only routing, no flow control or packet ordering
 - end-to-end argument has been driving the internet development



routing (1)

■ the purpose of routing protocols: to ensure that packets are sent the best possible way through a

(sub)net. but what is a good path?

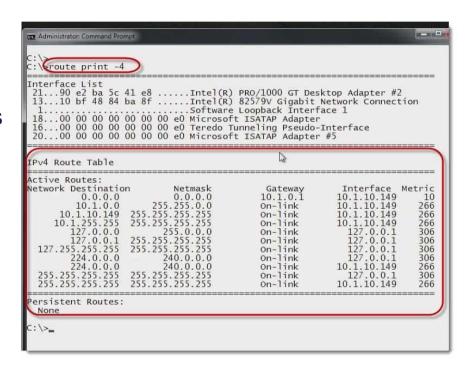


- several metrics:
 Hop count, Bandwidth, Delay, Jitter, Packet loss, Cost... Or combinations of the above...
- sometimes an advantage to use different paths through the network for the same session.
 - this can be part of a congestion control algorithm.
- in the Internet: packets are usually routed from router to router (hop by hop), where each router looks up the destination of the next.
 - o different packets can travel different paths...



routing (2)

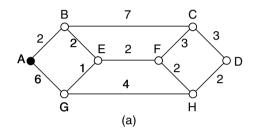
- handling routing tables and tables updated
- forwarding traffic according to routing tables

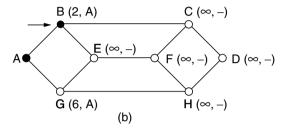


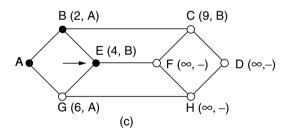


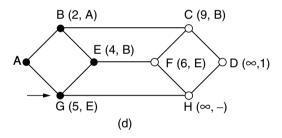
shortest path routing in static network

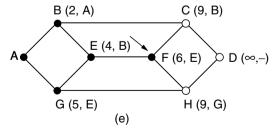
popular Dijkstra algorithm

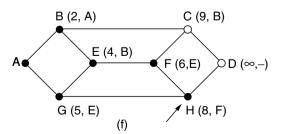














routing engine

- how can a router maintain an overview of a large network that can change dynamically?
- different methods and routing algorithms exist, such as:
 - flooding
 - Distance Vector Routing
 - Link State Routing
- flooding
 - replicate the incoming packet on all other links
 - many duplicates
 - introduce decreasing hop counter, ignoring of previously seen packets
 - only option when topology is unknown

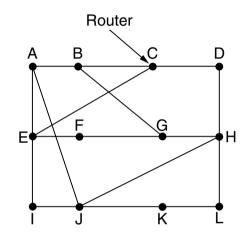


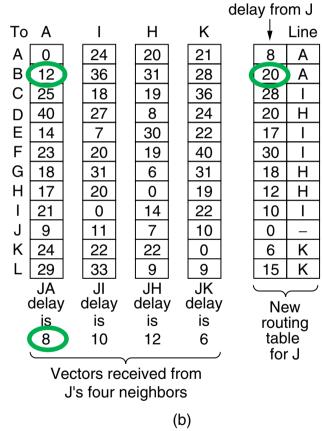
distance vector routing

- each node maintains routing table to the other routers in the network.
- each node stores the following information:
 - distance (measured in hops)
 - direction (what is the next hop)
- this information is spread to the neighbors of each node triggered by events or in fixed intervals
- upon receiving this information, it can re-calculate its routing tables.



distance vector routing: example





New estimated

AALBORG UNIVERSITY

(a)

distance vector algorithm (1)

based on Bellman-Ford algorithm

```
d_x(y) := \text{cost of least-cost path from } x \text{ to } y
then
d_x(y) = \min_{v} \left\{ c(x,v) + d_v(y) \right\}
cost \text{ from neighbor } v \text{ to destination } y
cost \text{ to neighbor } v
min \text{ taken over all neighbors } v \text{ of } x
```

distance vector algorithm (2)

key idea

- o from time-to-time, each node sends its own distance vector estimate to neighbours
- when a node receives an update from any neighbor, it updates its own DV using B-F equation
- convergence to the actual least cost to every destination in the network

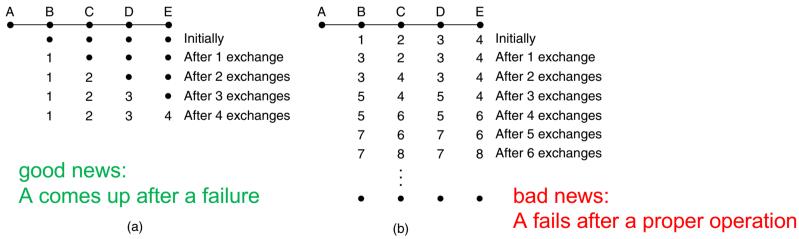
each node:

- → wait (change in local link cost or message update)
- → recompute
- → if gotten a new DV, notify neighbours
- → repeat the process
- completely distributed process



how long it takes to propagate through all network?

- with every algorithmic step, information is propagated one hop further
 - → time needed for convergence is linearly proportional to n
- "good news travel fast; bad news travel slow"
 - count-to-infinity problem



Petar Popovski, Communication in Electronic Systems, Fall 2024.

link cost changes

- convergence time varies (because it is iterative process)
- routing loops can occur (eventually, they will be resolved)
- security threaths: a router can advertise incorrect path cost → leading to black-holing
- reliability:
 unintentional error can propagate thru network while routing tables are updated

conclusion

- smart: very little overhead, since only communication with neighbours.
- not so smart: slow convergence, count-to-infinity problem.



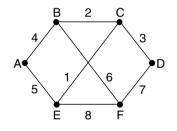
link state routing

- each node maintains a map of the network, from which distances can be calculated using standard shortest-path algorithms (e.g. Dijkstra).
- the maps are maintained by each node flooding information to all other nodes about itself and the distance to its neighbours.
- variants of link state routing that are widely used on the internet
 - IS-IS (Intermediate System to Intermediate System)
 - OSPF (Open Shortest Path First)



link state routing: steps

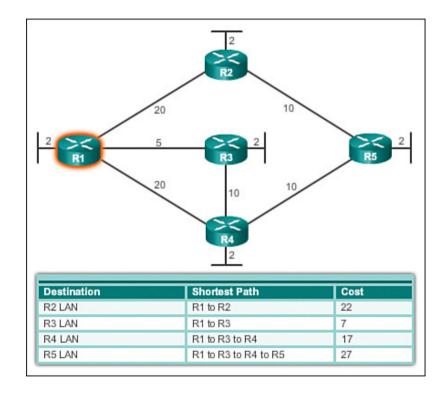
- discover its neighbors and learn their network addresses
 - HELLO messages
- set the distance or cost metric to each of its neighbors
 - inversely proportional to bandwidth of a link
- construct a packet telling all it has just learned.
 - o age: countdown for robustness
- send this packet to and receive packets from all other routers: flood
- compute the shortest path to every other router.



			Lir	ηk			S	ta	te		١	Pac	kets			
A	,	В)		D			Е	Ξ		F	=	
Seq.			Seq.			Seq.			Se	q.		Se	q.		Se	q.
Age			Age			Αç	ge		Αç	је		Αç	ge		Αç	ge
В	4		Α	4		В	2		О	3		Α	5		В	6
П	5		O	2		D 3		F 7			O	1		О	7	
			F	6		Е	1					F	8		Е	8

link state routing: features

- centralized:
 network topology and link costs are known to
 ALL nodes
 - Link State broadcast messages
 - all nodes will get the same info eventually
- Each node computes least cost paths from one node, itself, (source) to all other nodes → forwarding tables



Source: https://www.ciscopress.com/articles/article.asp?p=2180210&seqNum=11



complexity

- assume n nodes
- computation complexity
 - on iterations; for each iteration there is a comparison with nodes that are not selected yet $\rightarrow O(n^2)$ complexity
 - there exists a more efficient implementation that gives complexity of O(n logn)
- communication complexity
 - each node broadcasts its link state information to other n nodes
 - message complexity is $O(n^2)$

challenges and a verdict

- oscilations of link costs
 - especially, if costs related to available capacity or links congestion levels are used
- scalability
 - how to design the network so that it will still scale horizontally in 10 or more years but without becoming unmaintainable

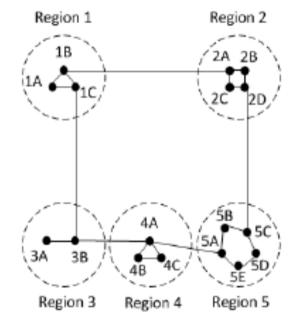
conclusions

- smart: no count-to-infinity problem
- not so smart: scalability

large networks: hierarchical routing

- large networks can be a challenge with respect to routing: table updates
 - table storage
 - look-up times

solution: hierarchical routing



other routing notions

- unicast (described so far)
- broadcast: send to all
- multicast: send to a group
- anycast: multiple nodes can have the response, so any response from them is fine

transport layer

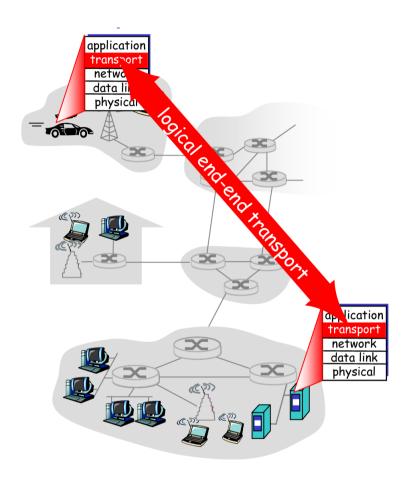




transport layer - functionality

- the transport layer is end-to-end addressing based on port numbers.
- it makes it possible to offer reliable and/or connection-oriented services over an unreliable and/or connectionless network.
 - o connection setup, maintenance and tear-down
 - retransmitting lost packets
 - handling out-of-order packet arrivals
- most used protocols:
 Transmission Control Protocol (TCP) and
 User Datagram Protocol (UDP).
- warning: simplifications may apply, a
 TCP is a huge subject,
 with many tweaks and optimization possibilities.

Petar Popovski, Communication in Electronic Systems, Fall 2024.



source:

http://www.sfu.ca/~ljilja/ENSC835/Spring09/News/Kurose_Ross/PowerPoint_Slides/Chapter3_4t h ed June 8 2007.pdf

some transport layer features and challenges

- goal: provide reliable transport service over a network of unreliable links
- code runs in all connected devices/machines; network-layer code in routers
- there can possibly be delays from the network layer
- retransmission of packets imply that there can exist multiple copies of the same data.
- packet loss is difficult to detect: when is a packet lost, and when it is delayed?
- it is often impossible to detect if a packet loss is due to congestion, or if it is real packet loss (on e.g. a wireless connection).



TCP

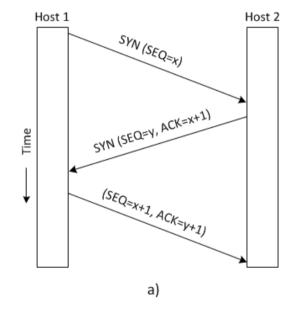
- it is defined in a number of RFCc: 793, 1122, 2018, 5681, 7323
- many versions
- byte stream, not a message stream
- TCP segment strcuture: consists of a *segment header* and a *data section*

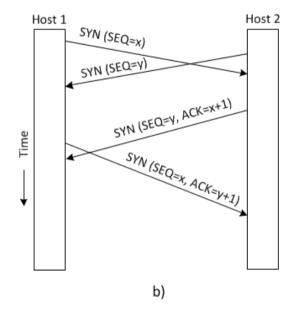
TCP segment header																																	
Offsets	Octet				0					2									3														
Octet	Bit	7 6 5 4 3 2 1 0 7 6 5 4 3 2 1 0 7 6 5 4 3										2	2 1	(0	7	6	5	4	3	2	1	0										
0	0	Source port														Destination port																	
4	32	Sequence number																															
8	64	Acknowledgment number (if ACK set)																															
12	96	Data offset Reserved 000 N S R E G K H T N N C R C S S Y I N N Window Size																															
16	128		Checksum Urgent pointer (if URG set)																														
20	160																																
:	:						Opti	ons	(if da	ata o	ffse	t > 5	. Pa	dded	l at t	he e	nd v	with	า "0)" b	its	if r	nece	ess	sary	/.)							
60	480																																



TCP connection set-up

- sequence numbers in one direction corresponds to acknowledgement numbers in the other.
 - o a) one connection.
 - b) one connection in each direction





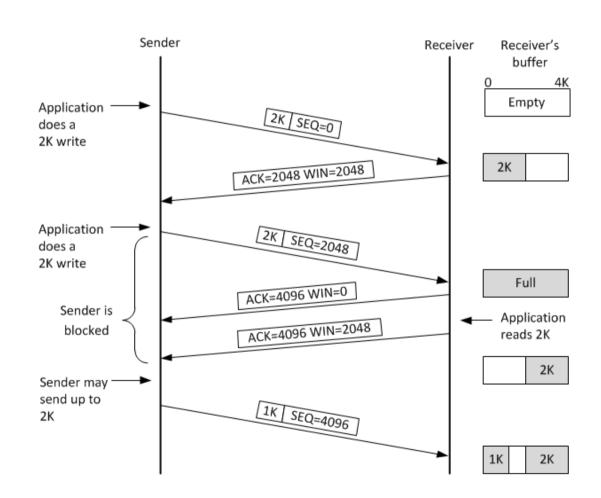
data transmission

how to achieve a reliable data transfer over an unreliable network?

- whenever data is sent, it is also acknowledged.
- in principle we could acknowledge every packet,
 but in TCP delayed acknowledgements can be used to reduce overhead
 - we wait a little before sending the acknowledgement,
 and if another packet arrives we can acknowledge them both in the same packet
 - o if a packet does not arrive, an acknowledgement is not sent, and it will be retransmitted.

TCP acknowledgements

- sliding windows
- ACK number tells
 which SEQ number to transmit next.
- if packets 1,2,3 are sent, but e.g. 2 is lost, the acknowledgement sent when packet 3 is received will be a duplicate ACK (and just the data from packet 1 will be acknowledged).
- once packet 2 is received,data can be acknowledged incl. packet 3.





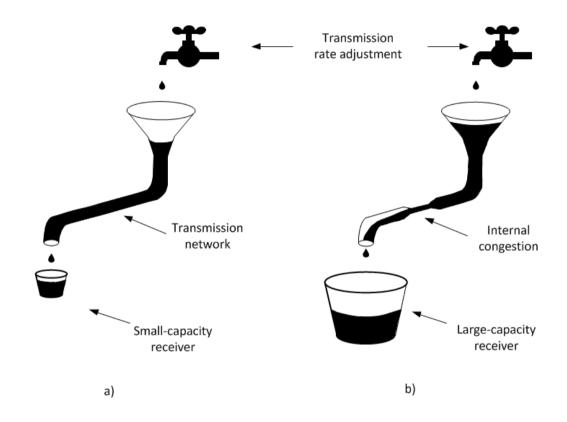
flow control

- there are two potential bottlenecks:
 - the network (or part of the network)
 - o the receiver

flow control:

preventing from overloading the receiver congestion control:

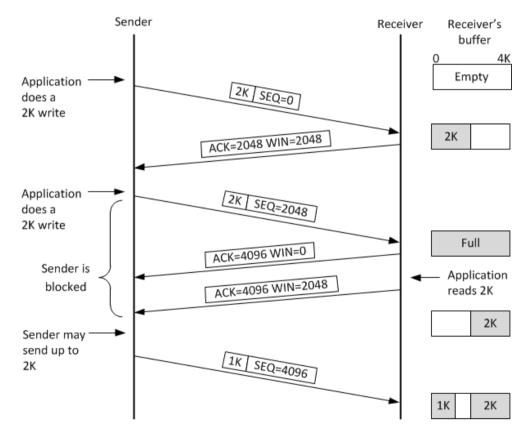
preventing from overloading the network





sliding windows: adjust to the receivers capacity

- the sender constantly maintains a picture of the receiver's window.
- it only sends data if the window is "open". ss soon as it is filled, the sender is blocked until an update is received.





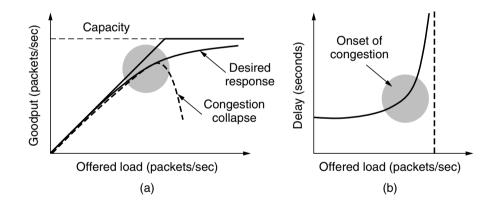
congestion control and fairness





congestion control

- what causes congestion?
 - o informally: "too many sources sending too much data too fast for network to handle"
- insights into congestion:
 - throughput can never exceed capacity
 - delay increases as capacity approached
 - loss/ retransmissions decreases the effective goodput
 - duplicates further decrease goodput



approaches towards congestion control

- end-to-end congestion control
 - o congestion window: the number of bytes the sender may have in the network at any time
 - o no explicit feedback from network (congestion "happens" on layer 3)
 - congestion is inferred from observed loss (and/or delay)
 - approach taken by TCP
- alternative approach = network assisted congestion control:
 - routers provide direct feedback
 - examples: TCP ECN (Explicit Congestion Notification), ATM

TCP congestion control: AIMD

- AIMD= additive increase, multiplicative decrease
- approach: senders can increase sending rate until packet loss (congestion) occurs, then decrease sending rate on loss event
- why AIMD?
 - o distributed, asynchronous algorithm
 - optimize congested flow rates network wide
 - have desired stability properties

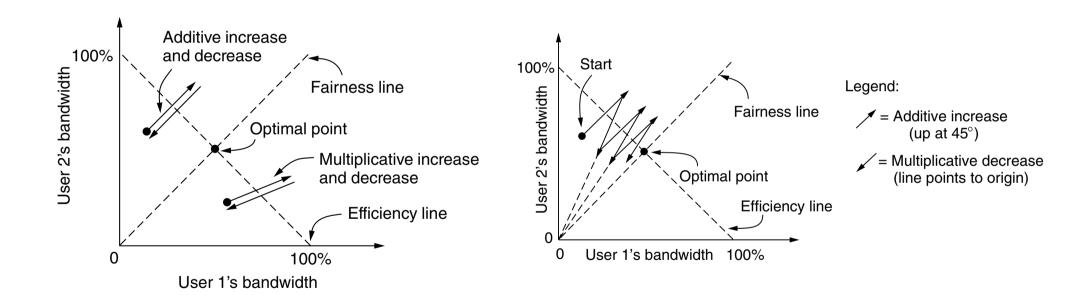


Time



why is AIMD fair

example: two competing data sessions



TCP congestion control: AIMD (contd.)

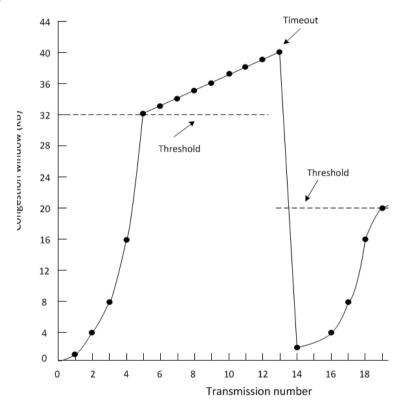
how much can be sent is regulated by the size of congestion window (cwnd)

slow start

- exponential increase of congestion window size
- doubling the amount of packets that we are sending

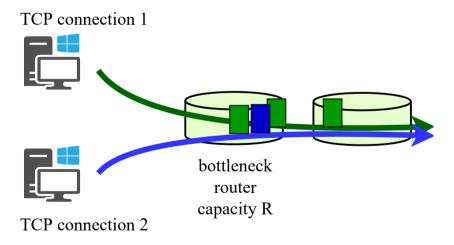
multiplicative decrease

- cut to 1 MSS (maximum segment size) when loss is detected by timeout (TCP Tahoe)
- cut in half on loss detected by triple duplicate ACK (TCP Reno)



TCP fairness

- fairness goal: if K TCP sessions share same bottleneck link of bandwidth R, each should have average rate of R/K
- max-min fairness: bandwidth of one flow cannot be increased without decreasing the one of another flow



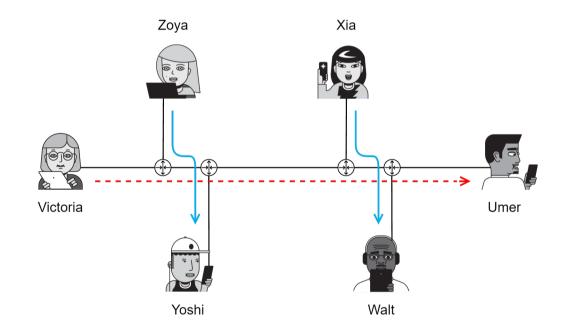


fairness

- multimedia apps often do not use TCP
- there is no "Internet police" policing use of congestion control
- no one prevents to use own unfair transport protocol,
 that will not "back off" in case of congestions
- parallel TCP connections
 - application can open multiple parallel connections between two hosts
 - they will be treated as separate TCP flows

a conflict between efficiency and fairness (1)

- 3 data flows
- Zoya-Yoshi and Xia-Walt flows
 - o high-resolution video call
 - saturate the horizontal links
- Victoria-Umer flow
 - audio call
 - lower data volume
- how to optimize the network?



a conflict between efficiency and fairness (2)

global efficiency

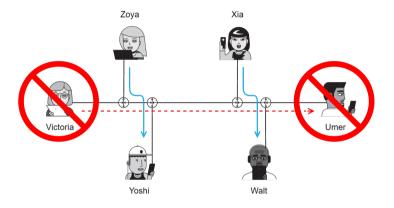
- shut down Victoria-Ulmer flow
- unfair network for Victoria and Ulmer

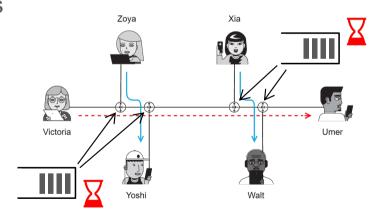
balancing efficiency and fairness

- routers queue and forward data packets
- increase packet delay for Zoya and Xia
- fair network for Victoria and Ulmer

trade-off control

- queueing and forwarding policies
- adaptive bitrate streaming

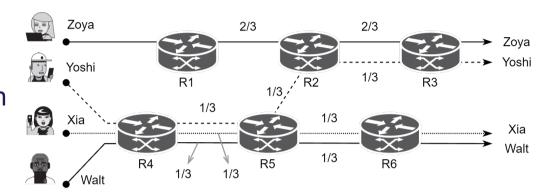






max-min fairness

- 4 data flows competing for bandwidth
- each link has a capacity equal to 1



- max-min fair bandwidth allocation
 - R4-R5: all flows get 1/3 of the bandwidth
 - R1-R2: Zoya's flow get 2/3 of the bandwidth
- why is the allocation max-min fair?
 - Yoshi's get 1/2 of the bandwidth at R2-R3
 - this increases Yoshi's bandwidth at R4-R5
 - o to implement this change, we need to decrease Xia's and Walt's bandwidth



summary and outlook

- networking layer and two main routing concepts
 - Distance Vector Routing
 - Link State Routing
- transport layer
 - basics of how TCP works and why it was designed in this way
 - connection and flow control
 - congestion control and fairness